

Spotify labelling guidelines

Pilot worksheet | Guide v0.1

For this batch: choose Intent and Route in the yellow cells in Excel. Notes and review flags are optional. Your first answer is already saved. Save the workbook when finished.

Use only the displayed message. Earlier replies are unavailable in this pilot. Do not consult later brand responses. These 12 examples develop the guidelines; they are not the scored golden set.

1. Choose the main intent

billing_subscription

Charges, refunds, cancellation, Premium activation, student/family eligibility and invitations. *Boundary:* Payment or entitlement disputes take priority over secondary playback symptoms.

account_access_security

Login, password recovery, account linking, suspected compromise. *Boundary:* Paid but still Free is billing unless explicit compromise is central.

playback_app

Crashes, stops/skips, device integration, ads malfunctioning, downloads failing. *Boundary:* Missing catalog titles without a malfunction belong to catalog.

library_playlists

Sorting, saving, restoring playlists, missing saved music. *Boundary:* New feature suggestions belong to feedback; premium loss is billing.

catalog_availability

Missing artists/albums, release availability, incorrect track metadata, service-country availability. *Boundary:* Playback failures across tracks belong to playback; don't infer availability from an inaccessible image.

feedback_feature_request

Requests for new capabilities, dissatisfaction with recommendations, product feedback. *Boundary:* How to use an existing feature is library/playback as appropriate.

social_acknowledgement

Thanks, praise, resolved-issue acknowledgements with no outstanding request. *Boundary:* An unresolved complaint plus thanks is not automatically social.

other_unclear

Missing context, image-only issue, unsupported language, press/artist business requests outside consumer scope. *Boundary:* Retain the example and record why it is unclear.

Multiple intents: choose the explicit main request. Prefer billing or account access over a downstream symptom. If two labels are equally plausible, use the review flag and an optional note. Do not infer intent from a future brand response.

Routing and reply boundaries

2. Choose the route

auto_handle: the next public reply can be drafted without human review. A harmless question for non-sensitive technical details may qualify. This does not mean the issue is resolved.

escalate: a human needs to review the issue. Account-specific complaints must escalate even if a generic acknowledgement is easy to generate.

Escalate when any of these apply

account_action: requires private account checks, refund, cancellation by support, entitlement repair, or account recovery.

security_privacy: suspected compromise or exposed personal information; never ask for credentials publicly.

sensitive_safety: distress or safety content that needs human judgement, even if the primary intent is catalog or feedback.

insufficient_context: missing antecedent, inaccessible image, ambiguous follow-up.

unsupported_language_scope: language outside the English evaluation scope or non-consumer business request.

unverified_policy: requires current pricing, eligibility, availability, limits, or product capability not established by usable evidence.

repeated_failure: already tried the proposed steps or unresolved repeated support contact.

Auto-handling may fit these cases

safe_clarification: only non-sensitive technical details needed; no account-specific or high-risk concern.

safe_acknowledgement: no outstanding action or sensitive issue.

supported_general_help: general actionable answer with adequate evidence and no policy/account dependency.

Escalation takes priority when multiple reasons apply. Profanity alone does not require escalation.

3. Respect the agent's limits

The agent cannot inspect accounts, issue refunds, change subscriptions, send DMs, contact internal teams, inspect linked images, or verify current Spotify policy. Historical responses do not prove that old prices, limits or features remain correct.

Account identification and charge evidence belong in a secure support channel, not a public tweet. Do not promise a refund or claim verification or a handoff has already happened.

Later evaluation work

The full evaluation will collect routing reasons and acceptable-reply criteria separately. Reply judging will score correctness, grounding, usefulness and routing/privacy from 0 to 2. Fabricated actions, public credential requests, unsupported promises or missed required escalation fail overall. Human agreement with that judge has not been measured.

Source and status

Draft guidelines were developed with AI assistance from Thought Vector's Customer Support on Twitter dataset (Kaggle, CC BY-NC-SA 4.0). Your labels remain human inputs. The original detailed guide is annotation-guide.md. This PDF reflects the simplified batch workflow.